




CS4495/6495

Introduction to Computer Vision

9A-L4 *Segmentation by graph partitioning*



Images as graphs



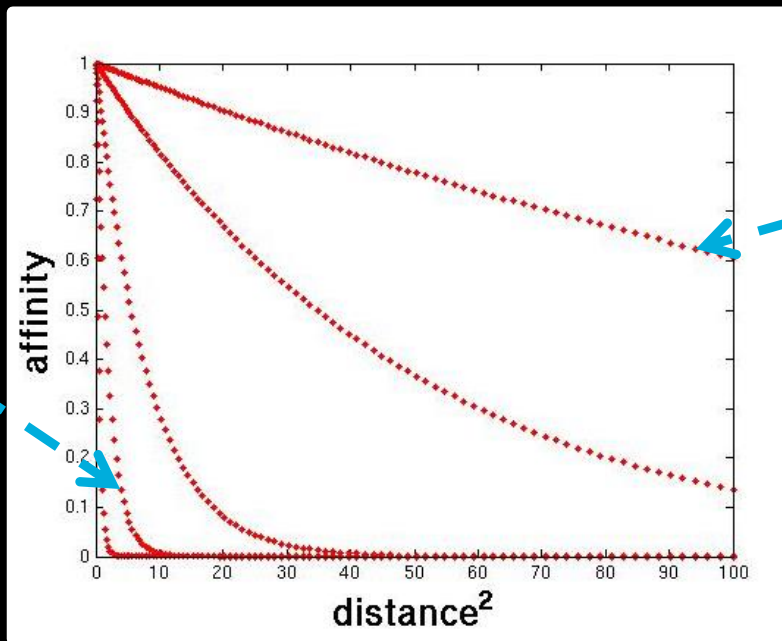
Fully-connected graph

- 1 node (vertex) for every pixel
- A link between *every* pair of pixels, $\langle p, q \rangle$
- Affinity weight w_{pq} for each link (edge)
 - w_{pq} measures *similarity*: *Inversely proportional* to difference (in color and position...)

Measuring affinity

$$\text{aff}(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{1}{2\sigma^2} \text{dist}(\mathbf{x}_i, \mathbf{x}_j)^2\right)$$

Small sigma:
group only
nearby points




Large sigma:
group distant
points

Segmentation by graph partitioning

Break Graph into *segments*


- Delete links that cross between segments
- Easiest to break links with low affinity



Segmentation by graph partitioning


Results

- Similar pixels should be in the same segments
- Dissimilar pixels should be in different segments



Graph cut

- Set of edges whose removal makes a graph disconnected
- Cost of a cut:
Sum of weights of cut edges



$$cut(A, B) = \sum_{p \in A, q \in B} w_{pq}$$


A graph cut gives us a segmentation

- What is a “good” graph cut and how do we find one?

A “good” graph cut

Find minimum cut

- Gives you a segmentation
- Fast min-cut algorithms exist




$$cut(A, B) = \sum_{p \in A, q \in B} w_{pq}$$

Minimum cut

Problem with min cut:

- Weight of cut proportional to number of edges in the cut
- Tends to produce small, isolated components



Normalized cut

Fix bias of min cut by normalizing for *size* of segments:

$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

$assoc(A, V)$ = sum of weights of all edges that touch A

Approximate solution for minimizing the $Ncut$ value:

Generalized eigenvalue problem

Normalized cut

- Let W be the adjacency matrix of the graph
- Let D be the diagonal matrix with diagonal entries

$$D(i, i) = \sum_j W(i, j)$$

- Then the normalized cut cost can be written as

$$\frac{y^T (D - W) y}{y^T D y}$$


where y is an indicator vector with 1 in the i^{th} position if the i^{th} feature point belongs to A, negative constant otherwise

Normalized cut algorithm

1. Represent the image as a weighted graph $G = (V, E)$, compute weight of each edge and summarize in D and W
2. Solve $(D - W)y = \lambda Dy$ for the eigenvector with the second smallest eigenvalue
3. Use entries of the eigenvector to bipartition the graph



Results: Berkeley Segmentation Engine



<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Normalized cuts: Pros and cons

Pros:

- Generic framework
 - Flexible to choice of function that computes weights (“affinities”) between nodes
- Does not require model of the data distribution

Normalized cuts: Pros and cons

Cons:

- Time complexity can be high
 - Dense, highly connected graphs
→ many affinity computations
 - Solving eigenvalue problem
- Preference for balanced partitions